

NatureNLP Project Summary: Versions v3 to v6

This document outlines the evolution of the NatureNLP project from version 3 to version 6, highlighting major improvements, evaluation metrics, and our trajectory towards version 7. It serves as a professional report to showcase project progress to professors, collaborators, or potential investors.

Evaluation Summary: v3 – v6 vs. Baselines

Model	BLEU ↑	Perplexity ↓	Loss ↓	Tokens/sec ↑
NatureNLP v3	0.12	45.2	4.95	45
NatureNLP v4	0.16	35.3	4.32	73
NatureNLP v5	0.21	30.4	4.01	125
NatureNLP v6	0.29	23.7	3.81	310
GPT-2	0.17	32.0	3.8	110
GPT-Neo-125M	0.21	28.5	3.5	220
Grok-1	0.26	25.2	3.1	310

Version Evolution and Key Improvements

Version 3

- Initial baseline with basic fine-tuning.
- Low BLEU, high perplexity, slow inference speed.

Version 4

- Added dataset preprocessing improvements.
- Introduced padding/token optimizations.
- Gained 62% speedup over v3.

Version 5

- Applied dynamic batching + improved loss function.
- Revived 'dead neurons' via selective reactivation.

- First comparison against GPT-2 and Grok-1.
- Strong BLEU gains and notable speed increase.

Version 6

- Introduced multi-task learning (IMDB, GSM8K, GitHub Code).
- Modular fine-tuning pipeline with lightweight fallback.
- Clean tokenization and lower perplexity across board.

Upcoming in Version 7

- Rewrite tokenizer + data loader to fully adopt Sunflower Hive Structure.
- Add sparse attention & oscillatory layers.
- Compare with newer models (Mixtral, DeepSeek, Yi-34B).
- Target BLEU > 0.3, loss < 3.5, tokens/sec > 200.